

# 2014-05-09-cython.sagews

May 9, 2014

## Contents

|          |   |          |
|----------|---|----------|
| <b>1</b> | <b>Math 480b Sage Course</b>                                    | <b>1</b> |
| 1.1      | Overview of Sage . . . . .                                      | 1        |
| 1.2      | May 9, 2014 . . . . .   | 1        |
| 1.3      | The C Level . . . . .   | 1        |
| 1.4      | Motivation . . . . .  | 1        |
| 1.5      | First benchmark: pi-digits . . . . .                            | 1        |
| 1.6      | Next benchmark: Mandelbrot . . . . .                            | 2        |
| 1.7      | Next benchmark: binary-trees . . . . .                          | 2        |
| 1.8      | ??? . . . . .   | 2        |
| 1.9      | Problem: variance of a list of floating point numbers . . . . . | 2        |
| 1.10     | Why it is called Cython . . . . .                               | 4        |

## 1 Math 480b Sage Course

### 1.1 Overview of Sage

### 1.2 May 9, 2014

Screencast: REMIND ME

Plan

- Questions
- Cython

### 1.3 The C Level

- Python is like floating around underwater you can easily float around, etc., but you move slowly.
- C is like flying around in the air you go much more quickly, but can also easily crash and burn
- Cython makes it so you can fly

## 1.4 Motivation

- Browse <http://benchmarksgame.alioth.debian.org/u64/benchmark.php?test=alllang=python3lang2=gccdata=u64> and compare speed of languages on number crunching tasks.

## 1.5 First benchmark: pi-digits

Lets look at and try out the pi digits benchmark. Here Python and C are pretty close!

But note the use of gmpy, i.e., all the hard work in the python implementation happens at the C level.

```
# sage is way faster... (probably a different algorithm!)
%time s = str(N(pi, digits=10000))
CPU time: 0.01 s, Wall time: 0.01 s
```

## 1.6 Next benchmark: Mandelbrot

Oh crap, Python sucks

<http://benchmarksgame.alioth.debian.org/u64/performance.php?test=mandelbrot>

Of all languages, C is the fastest. And Python is almost the worst, being about 86 times slower.

I looked into the Python code, and it does the computation in parallel using 64 processes, even on a single core computer. On a 1600x1600 case, just the time to start those processes and do nothing takes more time than the C program takes to do everything

## 1.7 Next benchmark: binary-trees

This one is also very depressing for Python.

<http://benchmarksgame.alioth.debian.org/u64/performance.php?test=binarytrees>

C is first and Python is nearly last, being over 50 times slower.

## 1.8 ???

- why?
- what can we do?
- but Python is so nice and fun.
- Extending Python with C/C++: <https://docs.python.org/2/extending/extending.html>
- Cython: basically makes extending Python with C/C++ way, way easier than the official approach

Lets take a very, very simple example in Python and speed it up using Cython.

## 1.9 Problem: variance of a list of floating point numbers

variance = the mean of the squares of the difference of each value from the mean =  $\sum \frac{1}{n}(x_i - \mu)^2$ .

```
# straightforward Python implementation
```

```
def var0(v):
    m = float(sum(v))/len(v)
    return sum([(x-m)**2 for x in v])/len(v) # use ** so this is \
    standard Python
```

```
v = [random() for _ in range(10)]
v
[0.8589381826391068, 0.8681592440262227, 0.6892358810838731, 0.5800153806503944,
0.6653829882189602, 0.18009338714825207, 0.9077079599252226, 0.8509199973502831,
0.9766987594487242, 0.7177646364529632]
```

```
var0(v)
0.04736519346871963
```

```
v = [random() for _ in range(10000)]
```

```
%timeit var0(v)
25 loops, best of 3: 8.63 ms per loop
```

```
# Straightforward Cython implementation
```

```
# (if the "show auto-generated" doesn't work properly for you -- with \
yellow -- restart your project server, since I just fixed a bug in \
this.)
```

```
%cython
def var1(v):
    m = float(sum(v))/len(v)
    return sum([(x-m)**2 for x in v])/len(v)
```

```
https://cloud.sagemath.com/blobs/_projects_74af30b7_ad25_4308_a02e_c71fcd84de6e__sage_temp_
compute19dc0_9366_dir_wbVbwx_a_pyx_0.html?uuid=691f0279-4c77-4c59-bbb6-f2e9c9691f1dShowauto-generatedcode
```

```
%timeit var1(v)
125 loops, best of 3: 3.08 ms per loop
```

```
8.63/3.08
2.80194805194805
```

```
# Declare some types
```

```
%cython
def var2(list v):
    cdef double m, x
    m = float(sum(v))/len(v)
    return sum([(x-m)**2 for x in v])/len(v)
```

```
https://cloud.sagemath.com/blobs/stdsage.html?uuid=f56adb5e-24d0-44df-94b5-b03867d712dcShowauto-generatedcode
```

```
%timeit var2(v)
```

625 loops, best of 3: 786 s per loop

8.63/.786  
10.9796437659033

```
# Don't call Python's sum function, but do the sum ourselves; also don't \
    use square which is potentially slow.
```

```
%cython
def var3(list v):
    cdef double m, x, s, n
    n = len(v)
    s = 0
    for x in v:
        s += x
    m = s/n
    s = 0
    for x in v:
        s += (x-m)*(x-m)
    return s/n
```

<https://cloud.sagemath.com/blobs/stdsage.html?uuid=55f5ede0-a2bb-4ed2-a5a1-a6c6b67955fbShowauto-generated>

```
%timeit var3(v)
625 loops, best of 3: 209 s per loop
```

8.63/.209  
41.2918660287081

```
# Next, make our own data type for a list of doubles
```

```
%cython
cdef class DoubleList:
    cdef double* v
    cdef int n
    def __init__(self, v):
        self.n = len(v)
        self.v = <double*> sage_malloc(sizeof(double)*self.n)
        cdef int i
        for i in range(self.n):
            self.v[i] = v[i]

    def variance(self):
        cdef double m, x, s
        cdef int i, n
        n = self.n
        s = 0
        for i in range(n):
            s += self.v[i]
        m = s/n
```

```

s = 0
for i in range(n):
    x = self.v[i]
    s += (x-m)*(x-m)
return s/n

```

[https://cloud.sagemath.com/blobs/\\_projects\\_74af30b7\\_ad25\\_4308\\_a02e\\_c71fcd84de6e\\_\\_sage\\_temp\\_compute19dc0\\_9366\\_dir\\_rD2Ets\\_a\\_pyx\\_0.html?uuid=66159b17-c255-42be-908c-f6c61681bd01Showauto-generatedcode](https://cloud.sagemath.com/blobs/_projects_74af30b7_ad25_4308_a02e_c71fcd84de6e__sage_temp_compute19dc0_9366_dir_rD2Ets_a_pyx_0.html?uuid=66159b17-c255-42be-908c-f6c61681bd01Showauto-generatedcode)

```
vd = DoubleList(v)
```

```
%timeit vd.variance()
```

625 loops, best of 3: 29.8 s per loop

```
# WIN
```

```
8.63/.0298
```

```
289.597315436242
```

## 1.10 Why it is called Cython

[http://www.mohawkradio.com/images/merch\\_images/mens/cython\\_01.jpg](http://www.mohawkradio.com/images/merch_images/mens/cython_01.jpg)